




Charting My Path: Journey of a GSoC 2023 Contributor to PostgreSQL with pg_statviz


Rajiv Harlalka



One Fine Morning.....



Deadlocks



\$ whoami

Rajiv Harlalka




- A student from India
- Member of **Kharagpur Open Source Society** (kossiitkgp.org), a student led club at IIT Kharagpur, India
- (New) Member to the **PostgreSQL** community
- Terminal Guy, Loves Nature and plays Squash.



\$ whatis

pg_statviz



- A minimalist extension and utility pair

\$ whatis

pg_statviz



- A minimalist extension and utility pair
- Time series analysis and visualization of PostgreSQL internal statistics

\$ **whatis**

pg_statviz



- A mini...

```
rajiv in @192.168.29.145 pg_statviz
56% x  whatis pg_statviz
pg_statviz: Your PostgreSQL buddy
```

OR



A Day in Life of:

Perform Operations

Collect Statistics

Create metrics out of
the collected statistics

Generate plots/graphs/histograms

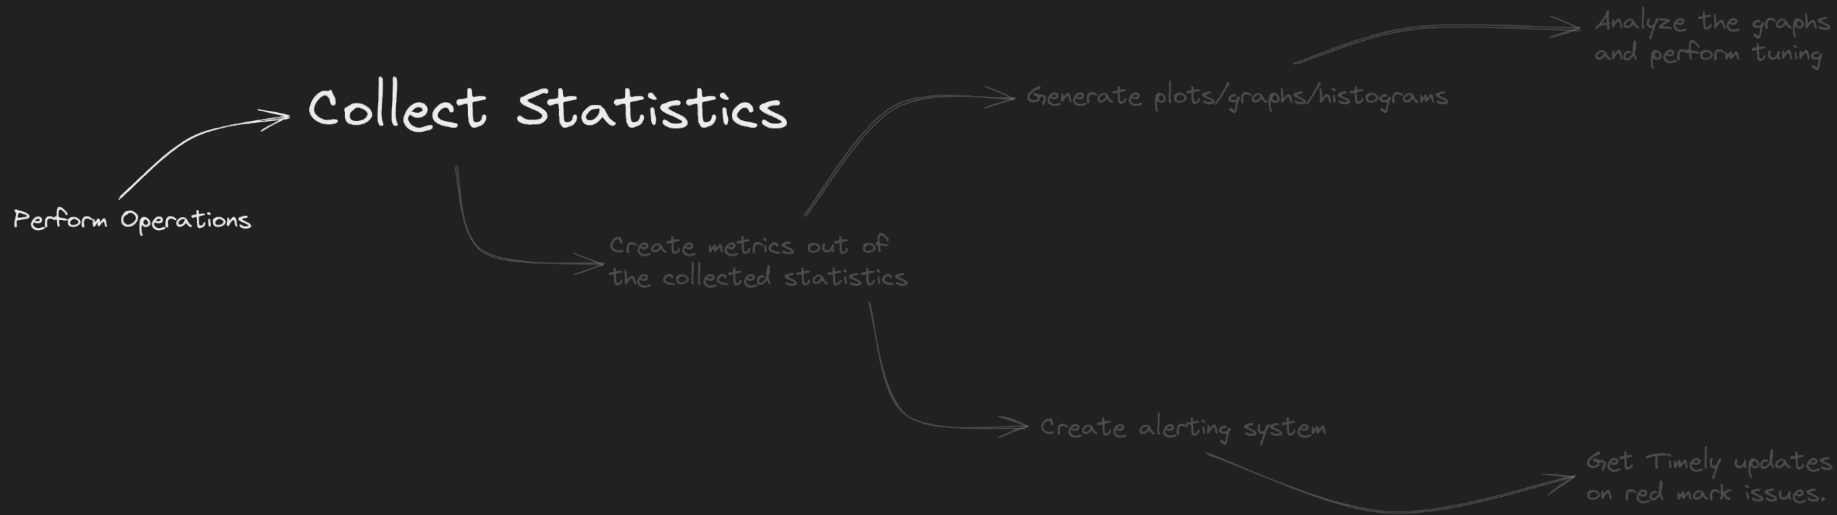
Analyze the graphs
and perform tuning

Create alerting system

Get Timely updates
on red mark issues.




A Day in Life of:



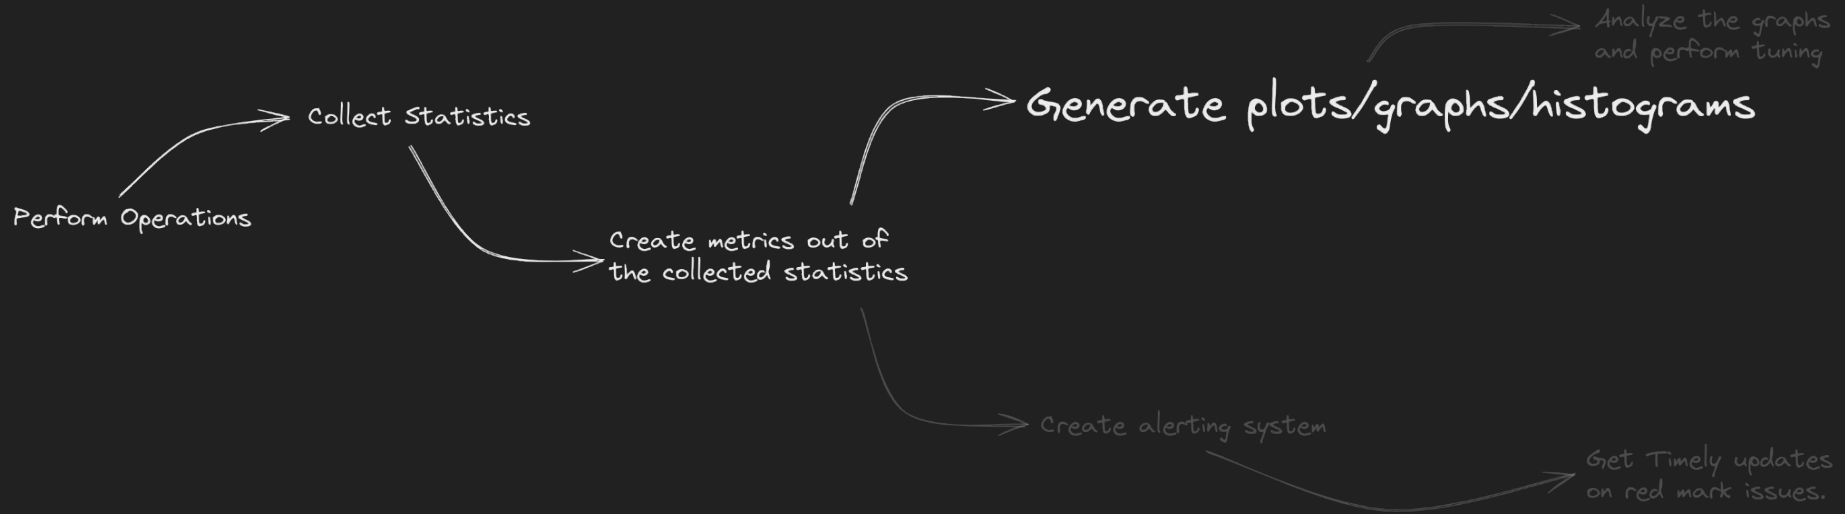


A Day in Life of:



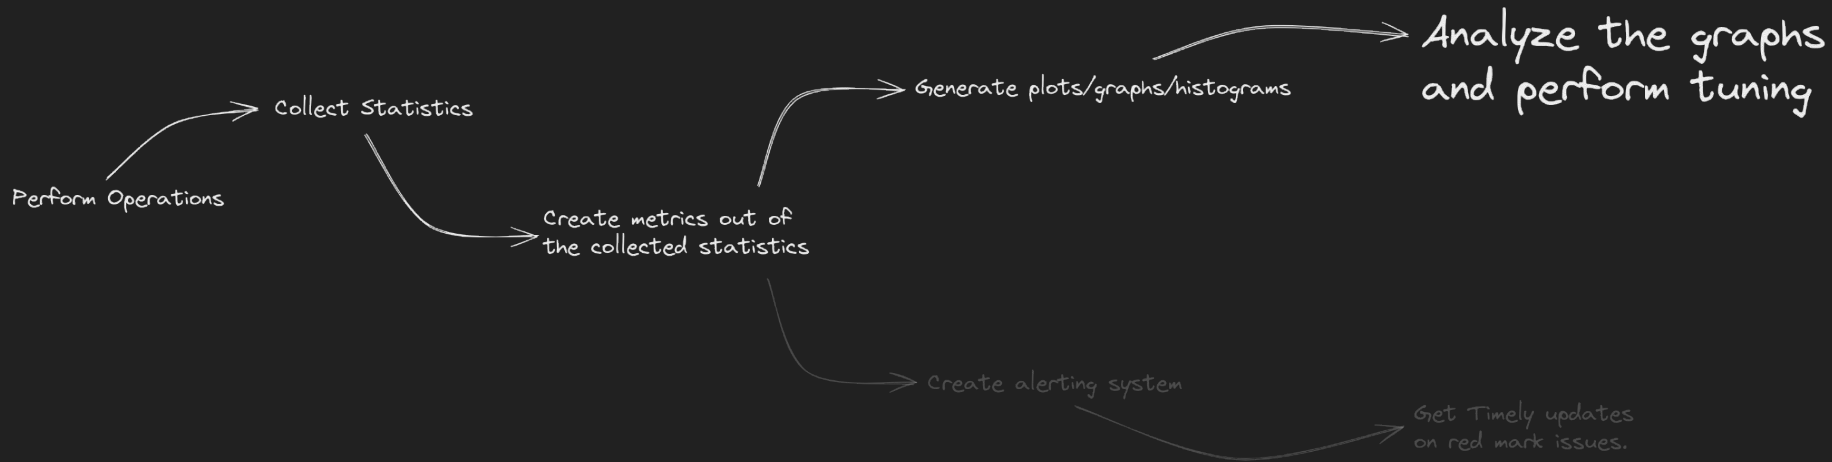


A Day in Life of:



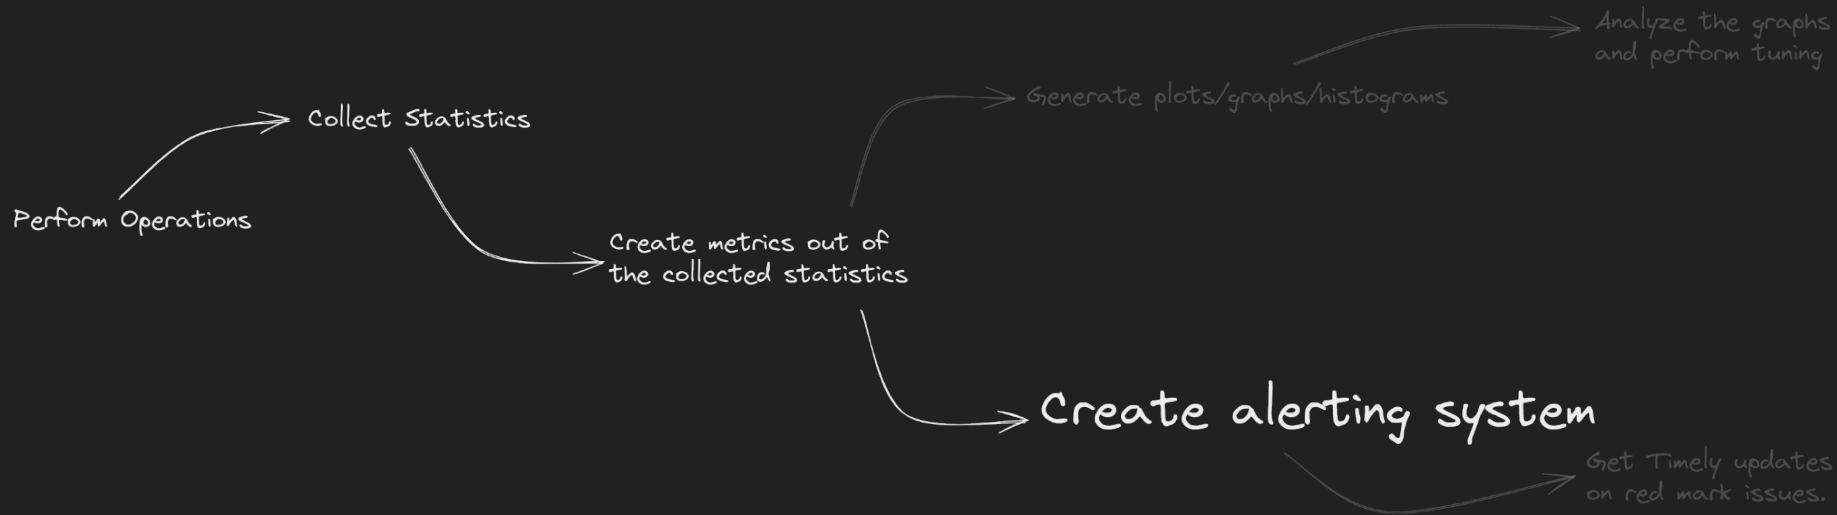


A Day in Life of:



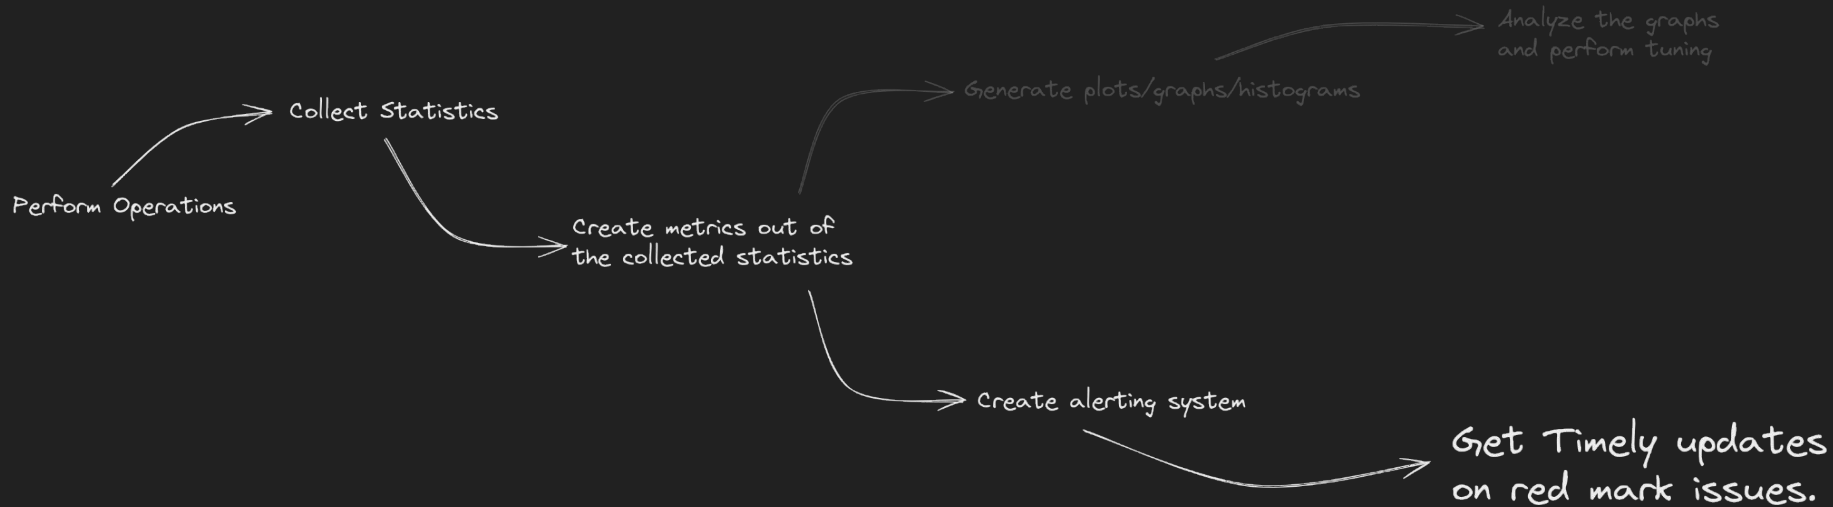


A Day in Life of:



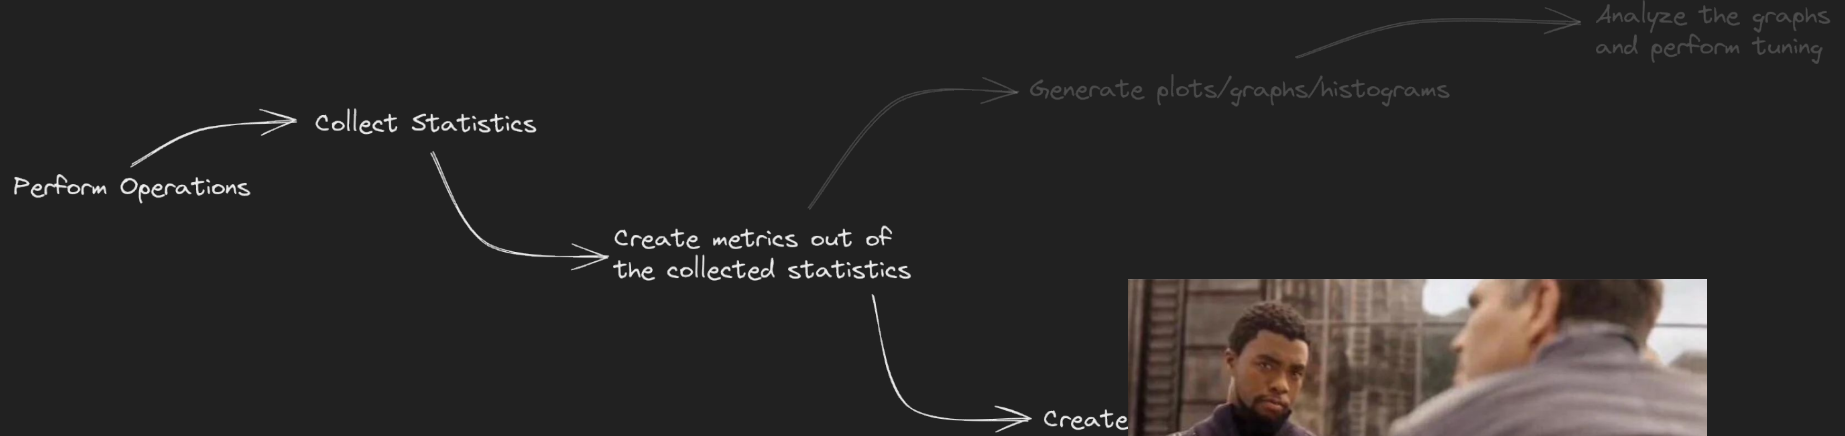


A Day in Life of:





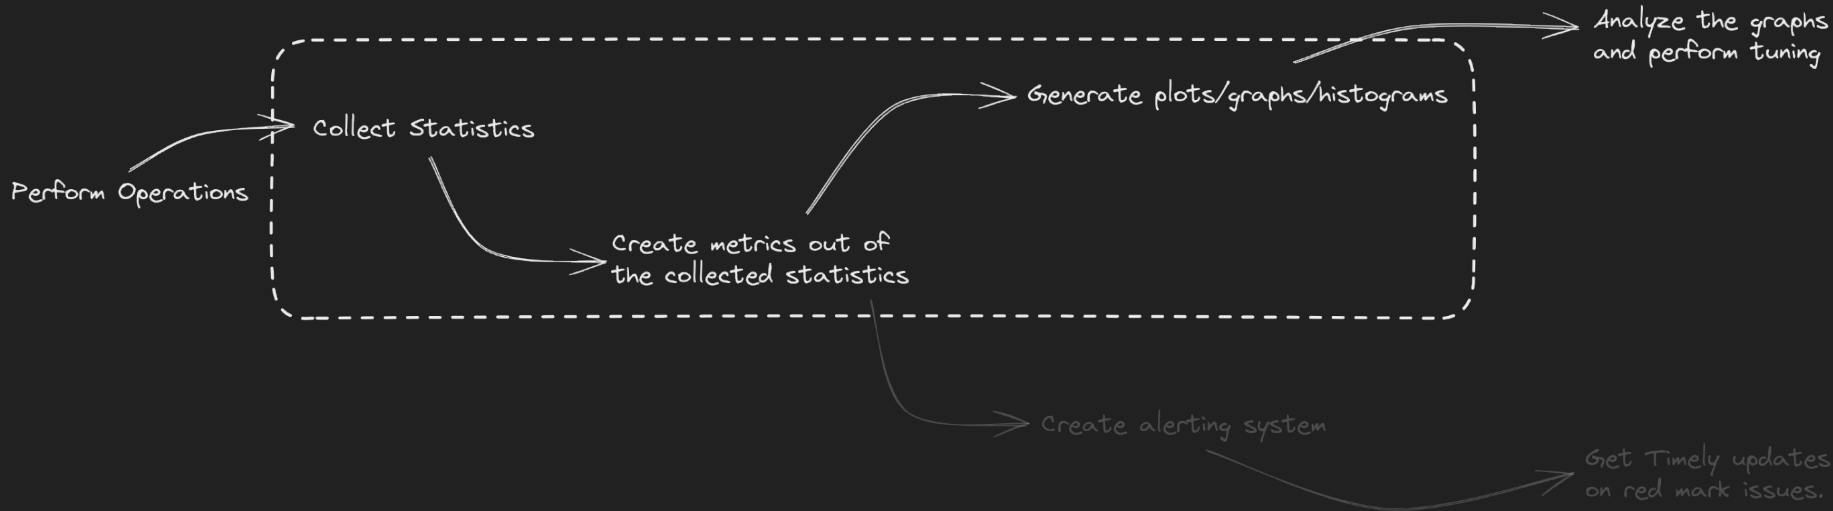
A Day in Life of:



y updates
rk issues.



A Day in Life of:



What pg_statviz focuses at



Why

pg_statviz

- Track PostgreSQL performance over time and potentially perform tuning or troubleshooting

Why

pg_statviz



- Track PostgreSQL performance over time and potentially perform tuning or troubleshooting

Why

pg_statviz



- Track PostgreSQL performance over time and potentially perform tuning or troubleshooting
- Understand about your system at a look.

Why

pg_statviz



- Track PostgreSQL performance over time and potentially perform tuning or troubleshooting
- Understand about your system at a look.
- Does not need knowledge about the internal tables of PostgreSQL.



Design Philosophy

- Minimal



Design Philosophy

- Minimal
- Modular



Design Philosophy


- Minimal
- Modular
- Designed with K.I.S.S and UNIX philosophy in mind



Design Philosophy

- Minimal
- Modular
- Designed with K.I.S.S and UNIX philosophy in mind

Does what it is meant to do.



Design Philosophy

- Minimal
- Modular
- Designed with K.I.S.S and UNIX philosophy in mind

Does what it is meant to do.

Provide Insights



Design Philosophy

Components:

- PostgreSQL Extension
- Python Utility Pair



Design Philosophy

Components:

- PostgreSQL Extension
- Python Utility Pair



Design Philosophy

Components:

- PostgreSQL Extension
- Python Utility Pair

No need to restart Database



Design Philosophy

Components:

- PostgreSQL Extension
- Python Utility Pair

No need to restart Database

No addition to `shared_preload_libraries`



Installation Guide

Getting hold of the extension:

- Get the packaged one

```
sudo dnf install pg_statviz_extension-<pg_version>
```

```
pgxn install pg_statviz
```



Installation Guide

Getting hold of the extension:

- Get the packaged one

```
sudo dnf install pg_statviz_extension-<pg_version>
```

```
pgxn install pg_statviz
```

- **Be Brave and build it:**
\$ **make** install



Installation Guide

Extension:

- **CREATE EXTENSION pg_statviz;**

Python Utility:

- **pip install pg_statviz;**



Installation

Extension:

- **CREATE**

Python Utility

- **pip inst**





Usage: Extension

- To create snapshots the user role should have either superuser or pg_monitor privilege



Usage: Extension

- Role should have either superuser or pg_monitor privilege

To take a snapshot:

```
SELECT pgstatviz.snapshot();
```



Usage: Extension

- Role should have either superuser or pg_monitor privilege

To take a snapshot:

SELECT pgstatviz.snapshot();

```
(postgres@localhost:5432) 12:12:12 [postgres]
# select pgstatviz.snapshot();
NOTICE:  created pg_statviz snapshot
          snapshot
-----
2023-12-05 12:12:20.228703+05:30
(1 row)
```



Usage: Utility Pair

```
68% → pg_statviz --help
```

```
usage: pg_statviz [--help] [--version] [-d DBNAME] [-h HOSTNAME] [-p PORT] [-U USERNAME] [-W] [-D FROM TO] [-O OUTPUTDIR]
               {analyze,buf,cache,checkp,conn,lock,tuple,wait,wal,xact} ...
```

run all analysis modules

positional arguments:

{analyze,buf,cache,checkp,conn,lock,tuple,wait,wal,xact}

analyze	run all analysis modules
buf	run buffers written analysis module
cache	run cache hit ratio analysis module
checkp	run checkpoint analysis module
conn	run connection count analysis module
lock	run locks analysis module
tuple	run tuple count analysis module
wait	run wait events analysis module
wal	run WAL generation analysis module
xact	run transaction count analysis module



Usage: Utility Pair

```
rajiv in @192.168.29.145 pg_statviz on 🌱 update-readme [📦👑] is 📦 v0.5
👹67% → ./src/run_pg_statviz -d postgres -h localhost -U postgres
INFO:pg_statviz.modules.buf:Running buffers written analysis
INFO:pg_statviz.modules.buf:Saving pg_statviz_rajiv_5432_buf.png
INFO:pg_statviz.modules.buf:Saving pg_statviz_rajiv_5432_buf_rate.png
INFO:pg_statviz.modules.checkp:Running checkpoint analysis
INFO:pg_statviz.modules.checkp:Saving pg_statviz_rajiv_5432_checkp.png
INFO:pg_statviz.modules.checkp:Saving pg_statviz_rajiv_5432_checkp_rate.png
INFO:pg_statviz.modules.cache:Running cache hit ratio analysis
INFO:pg_statviz.modules.cache:Saving pg_statviz_rajiv_5432_cache.png
INFO:pg_statviz.modules.conn:Running connection count analysis
INFO:pg_statviz.modules.conn:Saving pg_statviz_rajiv_5432_conn_status.png
INFO:pg_statviz.modules.conn:Saving pg_statviz_rajiv_5432_conn_user.png
INFO:pg_statviz.modules.lock:Running locks analysis
INFO:pg_statviz.modules.lock:Saving pg_statviz_rajiv_5432_lock.png
INFO:pg_statviz.modules.tuple:Running tuple count analysis
INFO:pg_statviz.modules.tuple:Saving pg_statviz_rajiv_5432_tuple.png
INFO:pg_statviz.modules.wait:Running wait events analysis
INFO:pg_statviz.modules.wait:Saving pg_statviz_rajiv_5432_wait.png
INFO:pg_statviz.modules.wal:Running WAL generation analysis
INFO:pg_statviz.modules.wal:Saving pg_statviz_rajiv_5432_wal.png
INFO:pg_statviz.modules.wal:Saving pg_statviz_rajiv_5432_wal_rate.png
```



PostgreSQL Internal Statistics

- Statistics on Internal activity collected via Collector reported through views.
 - Most on table/index information on row & disk block levels
- 24 Cumulative Statistics Views
 - Collected through Statistics Collector
 - collection and reporting of information about server activity



PostgreSQL Internal Statistics

- Statistics on Internal activity collected via Collector reported through views.
 - Most on table/index information on row & disk block levels
- 24 Cumulative Statistics Views
 - Collected through Statistics Collector
 - collection and reporting of information about server activity



PostgreSQL Internal Statistics

- 10 Dynamic Statistics Views
 - Deals with statistics of things currently happening.
 - Eg: replication, vacuum
- Though cumulative, gone once reset.



PostgreSQL Internal Statistics

- 10 Dynamic Statistics Views
 - Deals with statistics of things currently happening.
 - Eg: replication, vacuum
- Though cumulative, gone once reset.



A look behind the curtain

Base Schema



A look behind the curtain



```
CREATE TABLE IF NOT EXISTS @extschema@.snapshots(  
    snapshot_tstamp timestamptz PRIMARY KEY  
);
```

Base Schema



A look behind the curtain





A look behind the curtain

```
(postgres@localhost:5432) 19:18:24 [postgres]
```


```
# \df pgstatviz.*
```

List of functions				
Schema	Name	Result data type	Argument data types	Type
pgstatviz	delete_snapshots	void		func
pgstatviz	snapshot	timestamp with time zone		func
pgstatviz	snapshot_buf	void	snapshot_tstamp timestamp with time zone	func
pgstatviz	snapshot_conf	void	snapshot_tstamp timestamp with time zone	func
pgstatviz	snapshot_conn	void	snapshot_tstamp timestamp with time zone	func
pgstatviz	snapshot_db	void	snapshot_tstamp timestamp with time zone	func
pgstatviz	snapshot_lock	void	snapshot_tstamp timestamp with time zone	func
pgstatviz	snapshot_wait	void	snapshot_tstamp timestamp with time zone	func
pgstatviz	snapshot_wal	void	snapshot_tstamp timestamp with time zone	func

```
(9 rows)
```



A look behind the curtain





```
cur.execute("""SELECT tup_returned, tup_fetched, tup_inserted, tup_updated,
                    tup_deleted, snapshot_tstamp, stats_reset
                FROM pgstatviz.db
                WHERE snapshot_tstamp BETWEEN %s AND %s
                ORDER BY snapshot_tstamp""",
            (daterange[0], daterange[1]))
data = cur.fetchmany(MAX_RESULTS)
if not data:
    raise SystemExit("No pg_statviz snapshots found in this database")

tstamps = [t['snapshot_tstamp'] for t in data]
returned = [t['tup_returned'] for t in data]
fetched = [t['tup_fetched'] for t in data]
inserted = [t['tup_inserted'] for t in data]
updated = [t['tup_updated'] for t in data]
deleted = [t['tup_deleted'] for t in data]
```


A look



! Python
script


```
waitkinds = []
for w in wevents:
    for e in w:
        if 'wait_event' in e:
            wk = {'wait_event_type': e['wait_event_type'],
                  'wait_event': e['wait_event']}
            if wk not in waitkinds:
                waitkinds += wk,

# Plot as many of each wait event kind we have per snapshot
plt, fig = plot.setup()
plt.suptitle(f"pg_statviz · {info['hostname']}:{port}",
             fontweight='semibold')
plt.title("Wait events")
for wk in waitkinds:
    wc = []
    for w in wevents:
        if not w:
            wc += 0,
        else:
            found = False
            for e in w:
                if wk.items() <= e.items():
                    wc += e['wait_event_count'],
                    found = True
            if not found:
                wc += 0
```

! matplotlib
plotting !



A look behind the curtain





A look behind the curtain




```
def setup():
    for f in ["NotoSans-Regular.ttf", "NotoSans-SemiBold.ttf"]:
        f = importlib.resources.files("pg_statviz.libs").joinpath(f)
        fnt.fontManager.addfont(f)
    plt.rcParams['font.family'] = 'Noto Sans'
    plt.rcParams['font.size'] = 12
    base_image_path = importlib.resources.files("pg_statviz.libs")\
        .joinpath("pg_statviz.png")
    im = plt.imread(str(base_image_path))
    height = im.shape[0]
    fig = plt.figure(figsize=(19.2, 10.8))
    fig.figimage(im, 5, (fig.bbox.ymax - height - 6), zorder=3)
    plt.grid(visible=True)
    plt.ticklabel_format(axis='y', style='plain')
    plt.gcf().autofmt_xdate()
    return plt, fig
```

lib
pg



A look behind the curtain






A look behind the curtain

```
import sys
from argh import ArghParser
from pg_statviz.modules.analyze import analyze
from pg_statviz.modules.buf import buf
from pg_statviz.modules.cache import cache
from pg_statviz.modules.checkp import checkp
from pg_statviz.modules.conn import conn
from pg_statviz.modules.io import io
from pg_statviz.modules.lock import lock
from pg_statviz.modules.tuple import tuple
from pg_statviz.modules.wait import wait
from pg_statviz.modules.wal import wal
from pg_statviz.modules.xact import xact
```



Let's play detective:


Checkpoint




pg_statviz · rajiv:5432

Checkpoints

Requested
Timed




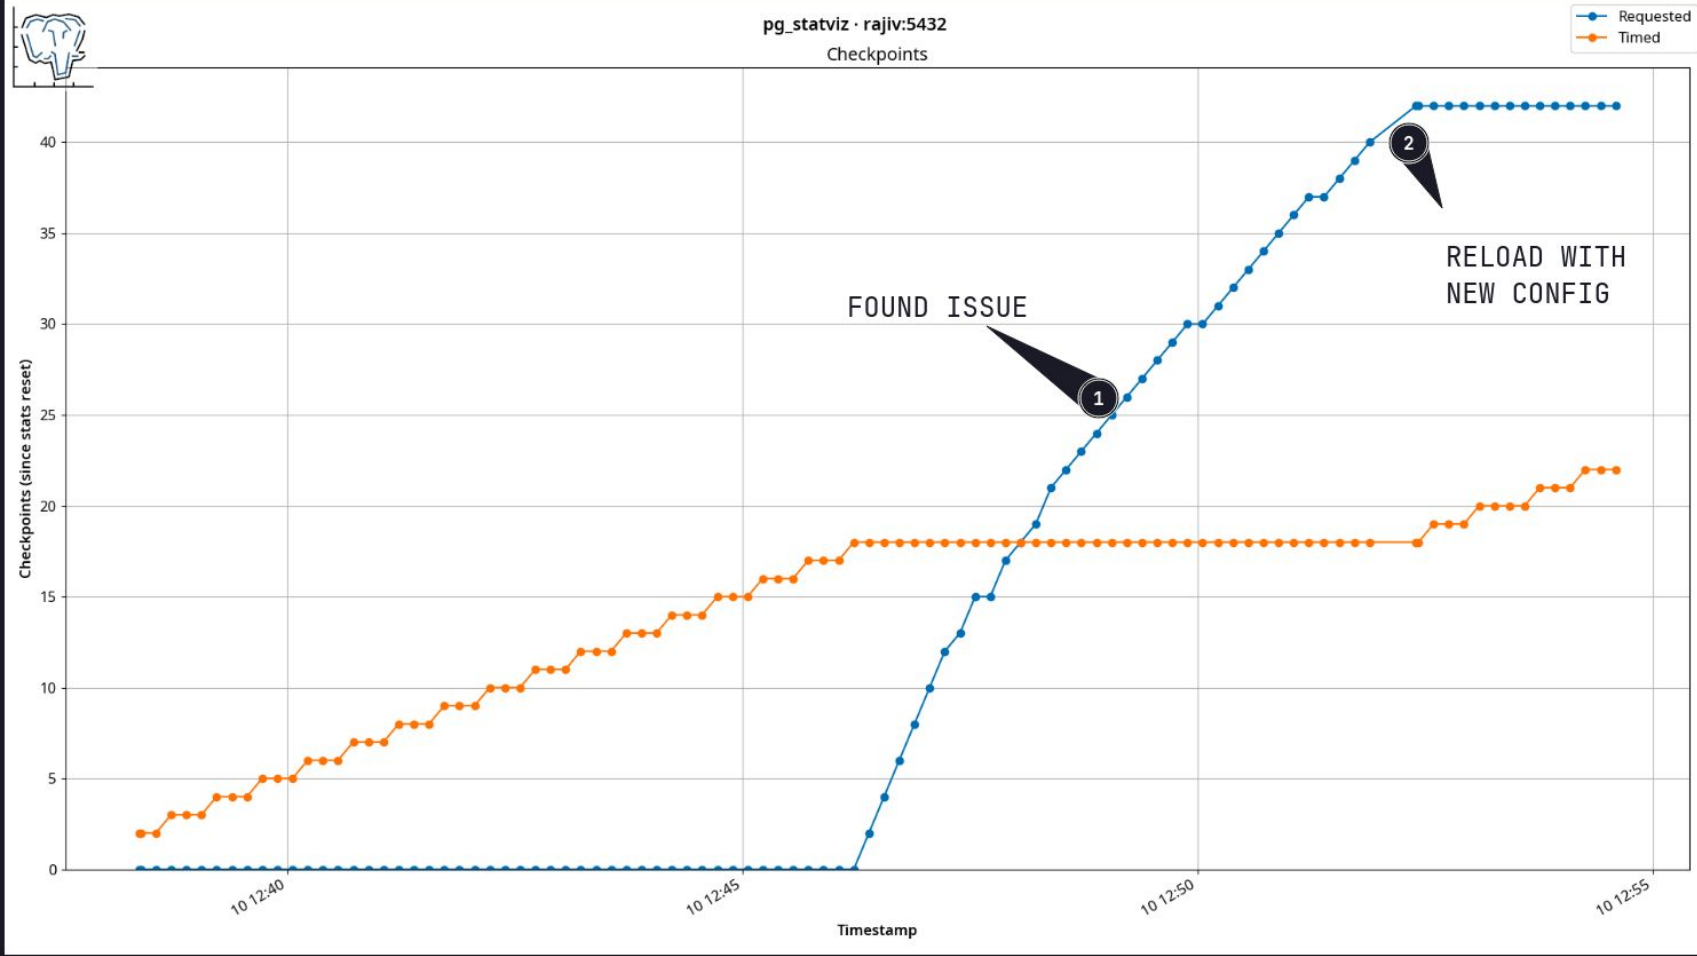



pg_statviz · rajiv:5432

Checkpoints

Requested
Timed








But what if **you** need a
Custom visualization

Create statistics collector



```
CREATE TABLE IF NOT EXISTS @extschema@.all_tab(  
    snapshot_tstamp timestamptz REFERENCES @extschema@.snapshots(snapshot_tstamp) ON DELETE  
    CASCADE PRIMARY KEY,  
    seq_scan int,  
    idx_scan int);  
  
CREATE OR REPLACE FUNCTION @extschema@.snapshot_all_tab(snapshot_tstamp timestamptz)  
RETURNS void  
AS $$  
    INSERT INTO @extschema@.all_tab (  
        snapshot_tstamp,  
        seq_scan,  
        idx_scan)  
    SELECT  
        snapshot_tstamp,  
        seq_scan,  
        idx_scan  
    FROM pg_stat_all_tables  
    WHERE relname = 'pgbench_accounts';  
$$ LANGUAGE SQL;
```

Create visualization generator



```
● ● ●

# Retrieve the snapshots from DB
cur = conn.cursor()
cur.execute("""SELECT snapshot_tstamp, seq_scan,
idx_scan

                FROM pgstatviz.all_tab
                WHERE snapshot_tstamp BETWEEN %s

AND %s

                ORDER BY snapshot_tstamp DESC""",
            (daterange[0], daterange[1]))
data = cur.fetchmany(MAX_RESULTS)
if not data:
    raise SystemExit("No pg_statviz snapshots
found in this database")

tstamps = [ts['snapshot_tstamp'] for ts in data]
seq_scan = [t['seq_scan'] for t in data]
idx_scan = [t['idx_scan'] for t in data]
```


```
● ● ●

# Plot as many of each scan type kind we have per
snapshot
plt, fig = plot.setup()
plt.suptitle(f"pg_statviz · {info['hostname']}:{port}",
            fontweight='semibold')
plt.title("Seq Scan vs Idx Scan on pgbench_accounts")
# Plot total wait events
plt.plot_date(tstamps, seq_scan,
              label='seq_scan', aa=True,
              linestyle='solid')
plt.plot_date(tstamps, idx_scan,
              label='idx_scan', aa=True,
              linestyle='solid')
fig.axes[0].set_ylim(bottom=0)

fig.gca().yaxis.set_major_locator(MaxNLocator(integer=True))
plt.xlabel("Timestamp", fontweight='semibold')
```



pg_statviz · rajiv:5432
Seq Scan vs Idx Scan on pgbench_accounts





Before we move on, a bit how did I
manage to do all this while
pursuing my college...



Google Summer of Code (GSoC)

- Global, online program focused on bringing new contributors into open source software development



Google Summer of Code (GSoC)

- Global, online program focused on bringing new contributors into open source software development
- Started in 2005 by Google and PostgreSQL has been a part since 2006



Google Summer of Code (GSoC)

- Global, online program focused on bringing new contributors into open source software development
- Started in 2005 by Google and PostgreSQL has been a part since 2006
- Contributors are paired with mentors from open source organizations, gaining exposure to real-world software development techniques



Opportunity for Students

- Chance to work on real world open source projects of interest
- Receive mentoring from experienced people in the industry
- Connect to a larger community of the org and fellow GSoCer's.
- Monetary and other Incentives



Opportunity for Students

- Chance to work on real world open source projects of interest



Opportunity for Students

- Chance to work on real world open source projects of interest
- Receive mentoring from experienced people in the industry



Opportunity for Students

- Chance to work on real world open source projects of interest
- Receive mentoring from experienced people in the industry
- Connect to a larger community of the org and fellow GSoCer's.



Opportunity for Students

- Chance to work on real world open source projects of interest
- Receive mentoring from experienced people in the industry
- Connect to a larger community of the org and fellow GSoCer's.
- Monetary and other Incentives



Opportunity for Orgs/Mentors

- Get contributions to your project



Opportunity for Orgs/Mentors

- Get contributions to your project
- Provide opportunity to students to begin their open source journey



Extending the tool



Extending the tool

- BYOM(Bring your own metric)



Extending the tool

- BYOM(Bring your own metric)
 - Snapshot the necessary statistics required.



Extending the tool

- BYOM(Bring your own metric)
 - Snapshot the necessary statistics required.
 - Update extension on the database



Extending the tool

- BYOM(Bring your own metric)
 - Snapshot the necessary statistics required.
 - Update extension on the database
 - Create module for your metric.



Extending the tool

- BYOM(Bring your own metric)
 - Snapshot the necessary statistics required.
 - Update extension on the database
 - Create module for your metric.
 - Export metrics in form of graphs.



Extending the tool

- BYOM(Bring your own metric)
 - Snapshot the necessary statistics required.
 - Update extension on the database
 - Create module for your metric.
 - Export metrics in form of graphs.
- Create live monitoring dashboards.



Extending the tool

- BYOM(Bring your own metric)
 - Snapshot the necessary statistics required.
 - Update extension on the database
 - Create module for your metric.
 - Export metrics in form of graphs.
- Create live monitoring dashboards.
- **An alerting solution**

Wrapping up..

- A simple tool to gain insights on your database
- Doesn't involve dependencies
- Not required to restart the db
- Modular
- Open to contributions



Thanks for tuning in...



[linkedin.com/in/rajivharlalka](https://www.linkedin.com/in/rajivharlalka)



github.com/rajivharlalka



twitter.com/rajivharlalka09

Project Link:

pg_statviz: github.com/vyruss/pg_statviz